


Smart System2: success vs. online Astra token cost

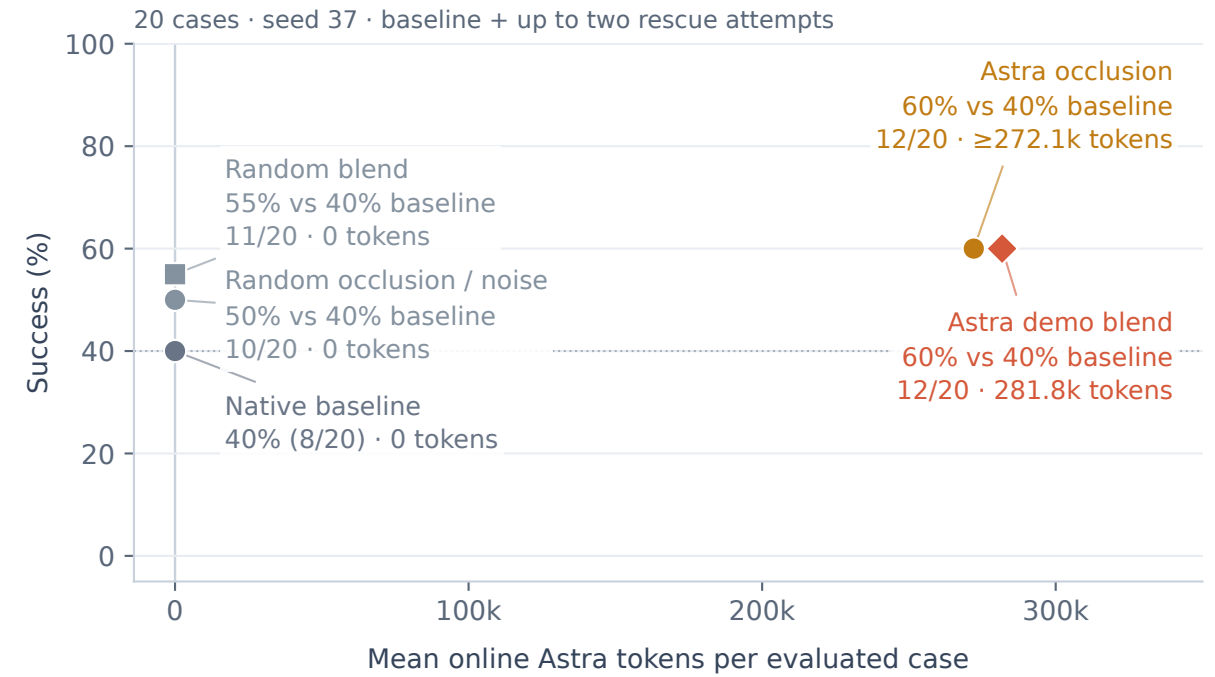
Higher success and fewer tokens are better. Separate cohorts; horizontal scales differ.

Cost = all reported input + output tokens / all evaluated task-reset cases, including failed attempts.

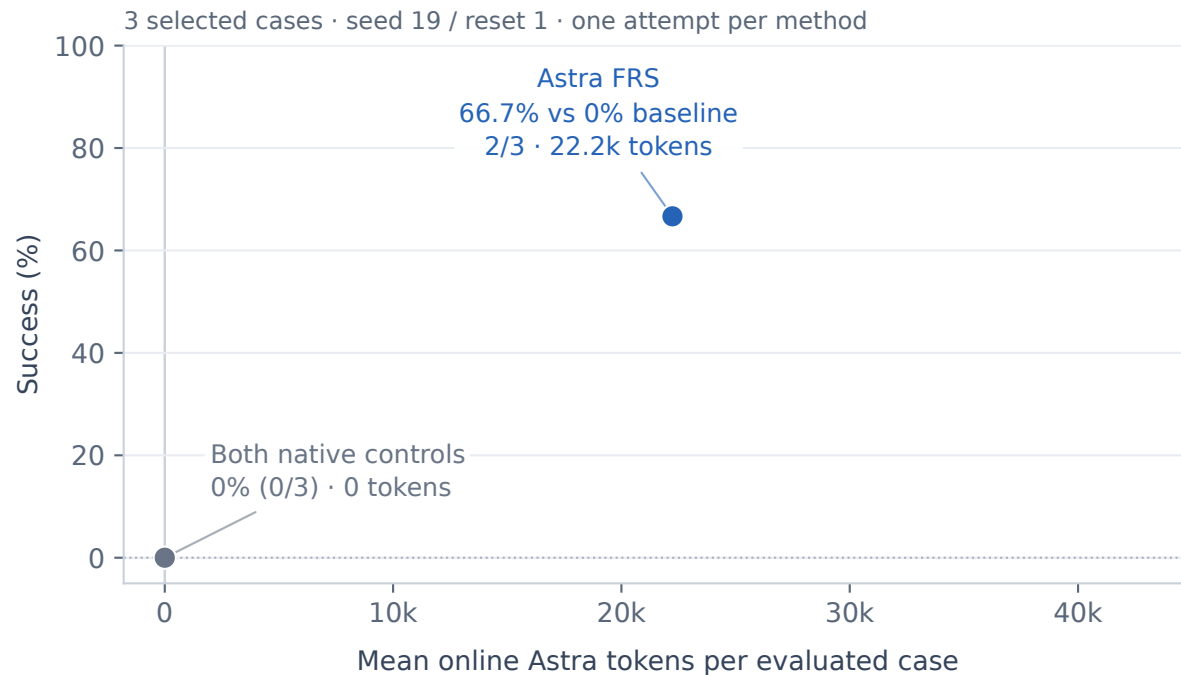
Experiment A · Language interventions



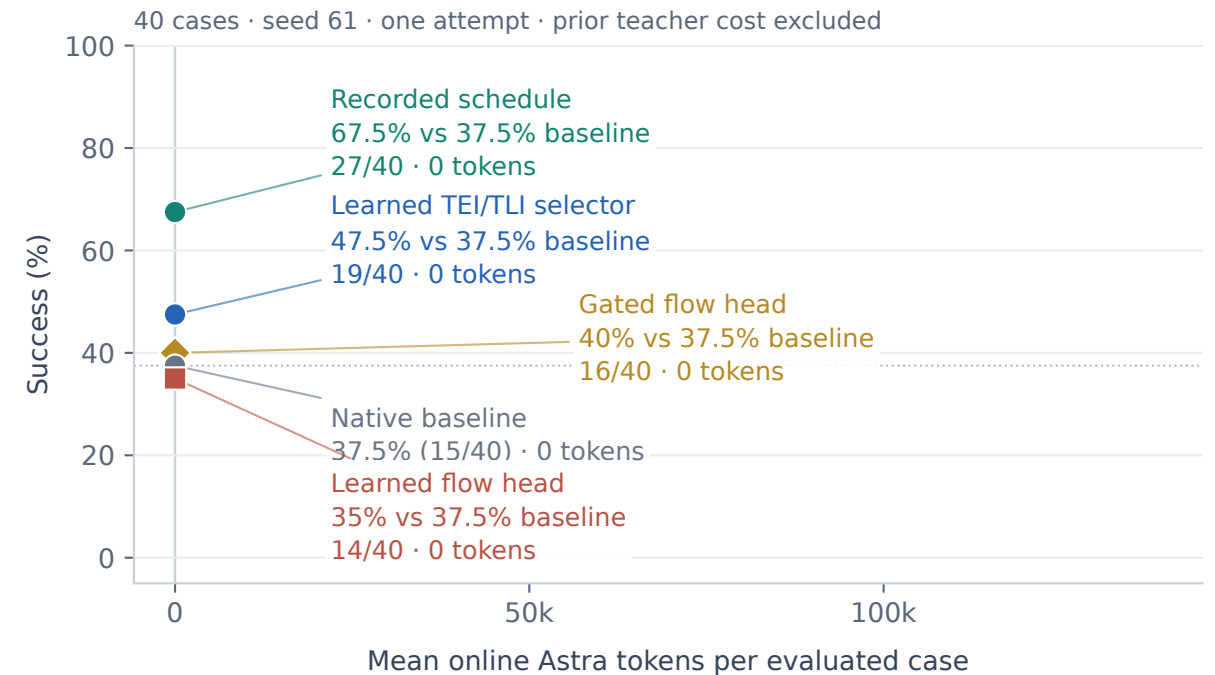
Experiment A · Image perturbations



Experiment A · FRS pilot



Experiment B · Reuse and distillation



≥ marks incomplete provider usage: one occlusion call has unknown usage, so its token cost is a lower bound.

Experiment B: zero online calls; whole teacher-source studies used ≥6.89M tokens. Exact selected-teacher acquisition cost is unknown.

FRS is a three-case development result. Compare methods within each panel; retries, resets and prompt/vision workloads differ.

This figure excludes development/interruption overhead for the 20-case studies, teacher acquisition, robot-policy compute and coding-agent tokens.

Larger FRS run: 952 recorded rollouts, 5/20 tasks fully audited before the spending cap; full-cohort SR is unavailable.